



KTH ROYAL INSTITUTE OF TECHNOLOGY

DD2424
DEEP LEARNING IN DATA SCIENCE

YEAR 2017 - 2018

Assignment 2

PICÓ ORISTRELL Sandra

940531 – 2308 sandrapo@kth.se

Professors :

SULLIVAN Josephine

August 5, 2018

Contents

1	Background	2
1.1	Background 1: Mathematical background.	2
1.2	Background 2: Speeding up training, add momentum.	3
1.3	Data structure	3
2	Exercises	4
2.1	Exercise 1: Read in the data and initialize the parameters of the network.	4
2.2	Exercise 2: Compute the gradients for the network parameters	4
2.3	Exercise 3: Add momentum to your update step	6
2.4	Exercise 4: Training your network	8
3	Conclusions	10
	Appendices	11
A	Assignment 2 implementation:	11

The main goal of this second assignment is to train and test a two-layer network with multiple outputs to classify images from the CIFAR-10 dataset. As in the first assignment, the network will be trained using mini-batch gradient descent with a cost function that computes the cross-entropy loss and with an L_2 regularization term on the weight matrix. The main structure of the code will be similar to assignment 1, but adding more parameters as now we are implementing one extra layer.

In order to explain and show the results obtained in a proper way I will divide the report into different subsections.

At the same time, in the appendix section, I will append all the implemented functions to be able to check the code ,if needed.

1 Background

The first section of this assignment will be focused on explaining the background needed for the assignment implementation.

1.1 Background 1: Mathematical background.

Following the same structure than in the first assignment, the network that we are going to implement will be a classifier where the output will correspond to a vector of probabilities. Then, the main structure of the network will be as follows:



Figure 1: Main structure of the implemented network.

Then, the predicted class will correspond to the output label with higher probability, i.e the *arg max* of the probability's vector.

In order to compute the output vector we will implement the following two-layer neural network:

$$s_1 = W_1 x + b_1 \quad (1)$$

$$h = \max(0, s_1) \quad (2)$$

$$s = W_2 h + b_2 \quad (3)$$

$$\mathbf{p} = \text{softmax}(s) = \frac{\exp(s)}{1^T \exp(s)} \quad (4)$$

This means that during the training phase we will learn the parameters W_1 , W_2 , b_1 and b_2 .

Then, defining the model as: $\Theta = [W_1, W_2, b_1, b_2]$, we will solve the following optimization problem to be able to set all the mentioned trainable parameters:

$$\Theta^* = \underset{\Theta}{\operatorname{argmin}} J(D, \lambda, \Theta) \quad (5)$$

Where J is a cost function that minimize the cross-entropy plus a regularization term on W .

$$J(D, \lambda, \Theta) = \frac{1}{|D|} \sum_{x,y \in D} \log(y^T p) + \lambda \sum_{l=1}^2 \sum_{i,j} W_{l,i,j}^2 \quad (6)$$

The problem will be solve via mini-batch gradient descent with momentum.

For the mini-batch implementation we will begin with sensible random initialization for W and b , and then estimate the parameters for $k = 1, 2$:

$$W_k^{(t+1)} = W_k^{(t)} - \eta \frac{\partial J^{(t+1)}, \lambda, \Theta}{\partial W_k} \quad (7)$$

$$b_k^{(t+1)} = b_k^{(t)} - \eta \frac{\partial J^{(t+1)}, \lambda, \Theta}{\partial b_k} \quad (8)$$

1.2 Background 2: Speeding up training, add momentum.

The vanilla version of the mini-batch gradient descent is really slow when using a sensible learning rate and with the size of data that we are going to use for this implementation. For this reason, we will need to speed up the training.

One of the possible ways to do it is adding a momentum term in the update step.

To achieve it, you first need to initialize a momentum vector (matrix) v_0 for each parameter of the network. The vector will be initialized to zero and will have the same dimension as the correspondence parameter. Once initialized, per each time step t :

$$v_t = p * v_{t-1} + \eta * \frac{\partial J}{\partial \theta} \quad (9)$$

$$\theta_t = \theta_{t-1} - v_t \quad (10)$$

where η is the learning rate, $p \in [0, 1]$ and θ is a generic placeholder to represent one of the parameters of the model.

In this assignment we will not implement an adaptive learning rate but we will decay the learning rate by some factor after each epoch.

1.3 Data structure

As in the previous assignment, the data that we are going to use is from the CIFAR-10 dataset. The dataset has a size of 10000×3072 following the next structure:




Figure 2: Data structure of CIFAR-10 dataset

2 Exercises

2.1 Exercise 1: Read in the data and initialize the parameters of the network.

To begin with this assignment we will split the data as follows:

- **Training data:** data_batch_1.mat
- **Validation data:** data_batch_2.mat
- **Test data:** test_batch.mat

To achieve it, we are going to use the function LoadBatch, already implemented for Assignment 1. Nevertheless, we will also need to pre-process the input data to have zero mean. We need to compute the mean for training data and then subtract this mean to the validation and test sets.

$$\text{mean_}X = \text{mean}(X, 2); \quad (11)$$

$$X = X - \text{repmat}(\text{mean_}X, [1, \text{size}(X, 2)]); \quad (12)$$

Next step will be initialize and set up all the parameters of the network. In this 2 layer neural network we will have $m=50$ nodes in the hidden layer. As the two weight matrices and bias vectors will have different dimensions, we will use cell arrays to store them.

The weight matrices will be initialized with a Gaussian distribution with 0 mean and 0.001 of standard deviation. On the other hand, bias vectors will be initialized to zero.

Then, the first assignment will be just the initialization of the network. The implementation can be found in the appendix section.

2.2 Exercise 2: Compute the gradients for the network parameters

After initialize and define all the network parameters, we will start writing the functions to compute the gradients. To do so, we will re-use the functions already implemented in Assignment 1.

Specifically, we are re-writing the function: ComputeGradients and transforming it to be able to compute a two-layer network using cell-arrays. EvaluateClassifier and ComputeCost are also modified.

Once the two mentioned functions are already implemented, it is important to check if the gradients are calculated properly. To be able to compare it, we will use the same methodology as in the first assignment; compare the gradients achieved with the numerical ones.

The achieved values are as follows:

	Values
$\mathbf{g_W_1}$	1.9894e-05
$\mathbf{g_b_1}$	7.1271e-04
$\mathbf{g_W_2}$	5.0976e-08
$\mathbf{g_b_2}$	6.2052e-10

As the values achieved are small enough, we will consider that the gradients are properly implemented.

The last experiment that we are going to implement to be sure that gradient computations and mini-batch gradient descent algorithm are okay is to train the network and check if it is possible to over-fit to training data and get a small loss after enough number of epochs.

Specifically, we are going to use the following configurations:

- Small amount of training data: 100 examples.
- Regularization term turned off: $\lambda = 0$.
- Enough number of epochs: 300.
- Reasonable learning rate: 0.01

As can be seen in the following plot, when we train the two layer network using the previous parameters, we are able to over-fit into the training data:

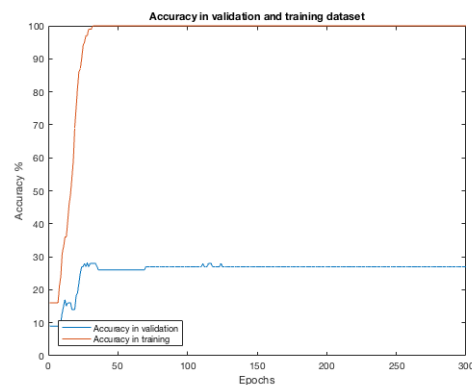


Figure 3: Training and testing accuracy with 100 examples.

We can conclude that the training is working because we are able to over-fit with only 29 epochs.

2.3 Exercise 3: Add momentum to your update step

During this third exercise, we have implemented the momentum vectors into the MiniBatch function to speed up the learning.

We will check its behaviour using different values of $\rho = [0.5, 0.9, 0.99]$ and fixing the learning rate into 0.001. The number of epochs will be set up to 10.

At the same time, to fully check the importance of it and the speed up achieved using momentum, I will do the following experiments using all the training data of the `data_batch_1.mat` dataset.

Here we can see the comparison between using momentum or without using it.

- **With momentum:**

- **$\rho = 0.5$**

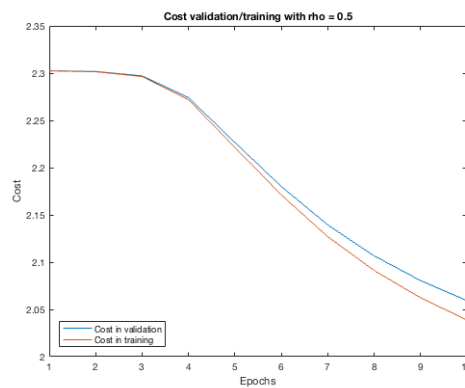


Figure 4: Cost with 10 epochs, $\rho = 0.5$

- **$\rho = 0.9$**




Figure 5: Cost with 10 epochs, $\rho = 0.9$

– $\rho = 0.99$




Figure 6: Cost with 10 epochs, $\rho = 0.99$

▪ **Without momentum:**

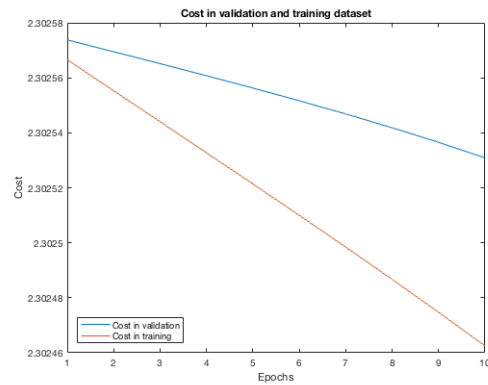


Figure 7: Cost with 10 epochs with no momentum.

As can be concluded from the previous images, using momentum helps to converge faster than not using it. Also, with momentum and $\rho = 0.90$ we achieve the best performance.

2.4 Exercise 4: Training your network

After doing the necessary experiments to verify that the momentum vectors and the gradients were well implemented, we will proceed to train the network correctly.

The first part of this assignment will be focused on finding a reasonable range of values for the learning rate parameter.

- **Find reasonable range of values for learning rate.**

During this exercise, we are going to use all the samples from the training data; i.e not only using 100 examples of the data.

As specified in the assignment description, we will use $\rho = 0.9$, a small value for the regularization term and also only 10 epochs.

To obtain a reasonable value of the learning rate, we will test all the next values:

`eta_values = [0.1,0.01,0.001,0.0001,1e-05,1e-06,1e-7];`

The obtained results can be seen in the next figure:




Figure 8: Training loss for different learning rate values.

As can be seen in the figure previously attached, the value of η greatly affects the loss obtained while training in the network. In the mentioned plot, we can conclude that when the value of the learning rate is too small; as in the cases of $1e-5$, $1e-6$ and $1e-7$, we practically do not learn; the loss is exactly the same during the first 10 epochs of the network.

Therefore, for the next experiment, we will consider that the first range of values of learning rate must be from 0.001 to 0.1.

- **Coarse random search.**

In this step we are going to find the best combinations for lambda and eta values, fixing $\rho = 0.9$.

The main intention will be performing training and then checking the learned network's best performance on the validation set; via the accuracy on it.

The first search will be focused over the feasible learning-rates that we identified in the previous experiment with a search over a very broad range of values for lambda.

Specifically, the range of values that we used for this first search are the next ones:

Network parameters	
rho	0.9
epochs	10
eta range values	0.001 to 0.1
lambda range values	1e-7 to 0.1
number of pairs	70

Using the mentioned configurations, the 3 best results obtained using this first search are:

Eta	Lambda	Validation accuracy
0.011778	0.000323	43.54%
0.013877	0.000165	43.28%
0.018709	0.007597	43.20%

- **Coarse-to-fine random search to set lambda and eta.**

During this search, we used a small range of values for lambda, to achieve better and accurate results. Specifically, we have been based on the previous best results. Then, during the coarse-to-fine random search, the parameters and the results obtained are as follows:

Network parameters	
rho	0.9
epochs	10
eta range values	0.01 to 0.07
lambda range values	1e-6 to 0.005
number of pairs	70

The results obtained in the coarse-to-fine random search are the next ones:

Eta	Lambda	Validation accuracy
0.014001	0.003347	44.00%
0.019811	0.003421	44.00%
0.018525	0.000001	43.82%

We can state that during this second search, we are improving the achieved results.

- **Best hyper-parameters**

Finally, using the best hyper-parameters combination found previously, we trained the model for longer and we checked the performance obtained on the test set.

During the training, we also break out early once the training cost is ever $> 3 \cdot$ original training cost.

Also, we are using only three of the training sets as training data instead of using all of them since exceeds Matlab maximum array dimensions.

Using the best hyper-parameters selection:

rho	lambda	eta
0.9	0.003347	0.014001

As a result of training the network during 30 epochs and with the previous specified hyper-parameters selection, we obtain 49.46% of accuracy in testing and 55.64% of accuracy in training data.

On the other hand, the validation and training loss achieved per each epoch, can be state in the figure:

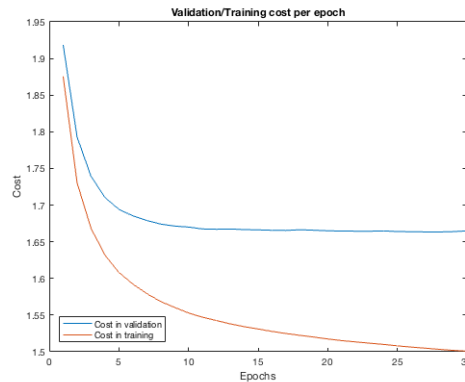


Figure 9: Final training and validation loss for 30 epochs.

3 Conclusions

During this assignment I learned about the implementation of a two-layer neural network, but above all the importance of correctly selecting the network hyper-parameters to be able to train and get better performance.

It should also be noted that without momentum, the experiments would have been much slower.

Therefore, with all the results obtained during the different experiments, we can conclude that the implementation of the 2-Layer neural network is properly implemented and verified.

Appendices

A Assignment 2 implementation:

```
1 %Author: Sandra Pic Oristrell
2 %Data: 13 of August 2018.
3 %Assignment 2: 2 layer neural network.
4
5 %% Exercise 1: Read in the data and initialize the parameters of the network.
6 %Load the data using the function implemented in assignment 1.
7
8 [X_training, Y_training, y_training] = LoadBatch('./data_batch_1.mat');
9 [X_validation, Y_validation, y_validation] = LoadBatch('./data_batch_2.mat');
10 [X_test, Y_test, y_test] = LoadBatch('./test_batch.mat');
11
12 %Apply more pre-processing to the raw input data.
13 %Transform it to have zero mean.
14 mean_X = mean(X_training, 2);
15 X_training = X_training - repmat(mean_X, [1, size(X_training, 2)]);
16 X_validation = X_validation - repmat(mean_X, [1, size(X_validation, 2)]);
17 X_test = X_test - repmat(mean_X, [1, size(X_test, 2)]);
18
19 %Variables
20 d = size(X_training,1);
21 N = size(X_training,2);
22 K = size(Y_training,1);
23
24 %Data structure for the parameters of the network and initialize the
25 %values.
26
27 %50 nodes in the hidden layer
28 m = 50;
29 lambda = 0;
30 %Initialize the network:
31 [W,b] = Initialize_Network(m,d,K);
32
33 %% Exercise 2: Compute the gradients for the network parameters.
34
35 %Re-write or update gradients functions from assignment 1.
36 %% 2.1 Check that the gradients are implemented properly.
37
38 batch_size = 100;
39
40 [P,h] = EvaluateClassifier(X_training(:,1:batch_size),W,b);
41 [grad_W, grad_b] = ...
    ComputeGradients(X_training(:,1:batch_size),Y_training(:,1:batch_size),P,h,W,b,lambda);
42 [ngrad_W, ngrad_b] = ...
    ComputeGradsNumSlow(X_training(:,1:batch_size),Y_training(:,1:batch_size),W,b,lambda,1e-5);
43
44 %Check the comparison between the numerical and the computed gradients:
45 eps = 1e-10;
46 %Layer 1:
47 gradient_b1_comparison = sum(abs(ngrad_b{1} - grad_b{1})/max(eps, sum(abs(ngrad_b{1}) ...
    + abs(grad_b{1})))));
48 gradient_W1_comparison = sum(sum(abs(ngrad_W{1} - grad_W{1})/max(eps, ...
    sum(sum(abs(ngrad_W{1}) + abs(grad_W{1}))))));
49 %Layer 2:
50 gradient_b2_comparison = sum(abs(ngrad_b{2} - grad_b{2})/max(eps, sum(abs(ngrad_b{2}) ...
    + abs(grad_b{2})))));
51 gradient_W2_comparison = sum(sum(abs(ngrad_W{2} - grad_W{2})/max(eps, ...
    sum(sum(abs(ngrad_W{2}) + abs(grad_W{2}))))));
```

```

52
53 %Check that the gradients have small number.
54 fprintf("Results for the calculated gradients:");
55 fprintf("Layer 1:");
56 fprintf("W1: %f",gradient_W1_comparison);
57 fprintf("b1: %f",gradient_b1_comparison);
58 fprintf("Layer 2:");
59 fprintf("W2: %f",gradient_W2_comparison);
60 fprintf("b2: %f",gradient_b2_comparison);
61
62 %% 2.2 Try if you can overfit in training data with 200 epochs and reasonable learning ...
    rate.
63
64 GDparams.eta = 0.01;
65 GDparams.n_batch = 1;
66 GDparams.n_epochs = 300;
67 lambda = 0;
68
69 %Save all the information per each epoch.
70 cost_training_list = zeros(1, GDparams.n_epochs);
71 cost_validation_list = zeros(1, GDparams.n_epochs);
72 accuracy_training_list = zeros(1, GDparams.n_epochs);
73 accuracy_validation_list = zeros(1, GDparams.n_epochs);
74 epochs_list = zeros(1, GDparams.n_epochs);
75
76 %Try with only 100 examples.
77 examples = 100;
78 X_train = X_training(:,1:examples);
79 Y_train = Y_training(:,1:examples);
80 y_train = y_training(:,1:examples);
81
82 X_val = X_validation(:,1:examples);
83 Y_val = Y_validation(:,1:examples);
84 y_val = y_validation(:,1:examples);
85
86 for i = 1:GDparams.n_epochs
87     [W,b] = MiniBatchGD(X_train, Y_train,y_train,GDparams,W,b,lambda);
88     fprintf("Epoch: %d\n",i);
89     epochs_list(i) = i;
90     cost_training_list(i) = ComputeCost(X_train, Y_train, W, b, lambda);
91     accuracy_training_list(i) = ComputeAccuracy(X_train,y_train,W,b);
92     accuracy_validation_list(i) = ComputeAccuracy(X_val,y_val,W,b);
93     cost_validation_list(i) = ComputeCost(X_val, Y_val, W, b,lambda);
94 end
95
96 %Plot cost in training and validation per epoch.
97 title_text = "Cost in validation and training dataset";
98 PlotCost(epochs_list, cost_validation_list, cost_training_list,title_text);
99
100 %Plot accuracy in training and validation per epoch.
101 title_text = "Accuracy in validation and training dataset";
102 PlotAccuracy(epochs_list,accuracy_validation_list, accuracy_training_list,title_text);
103
104 %Final accuracy in validation dataset and training dataset:
105 accuracy_final_training = ComputeAccuracy(X_train,y_train,W,b);
106 fprintf("Final training accuracy: %f %", accuracy_final_training);
107
108 accuracy_final_validation = ComputeAccuracy(X_val,y_val,W,b);
109 fprintf("Final validation accuracy: %f %",accuracy_final_validation);
110
111 %% Exercise 3: Add momentum to your update step.
112 %To help speed up training times you should add momentum terms into your mini-batch ...
    update steps
113
114 GDparams.eta = 0.001;

```

```

115 GDparams.n_batch = 100;
116 GDparams.n_epochs = 10;
117 lambda = 0;
118
119 eta_decay_rate = 0.95;
120
121 %Compare rho values = [0.5,0.9,0.99];
122 rho_values = [0.50,0.9,0.99];
123
124 for rho = rho_values
125     [W,b] = Initialize_Network(m,d,K);
126     GDparams.eta = 0.01;
127     cost_training_list = zeros(1, GDparams.n_epochs);
128     cost_validation_list = zeros(1, GDparams.n_epochs);
129     accuracy_training_list = zeros(1, GDparams.n_epochs);
130     accuracy_validation_list = zeros(1, GDparams.n_epochs);
131     epochs_list = zeros(1, GDparams.n_epochs);
132     for i = 1:GDparams.n_epochs
133         [W,b] = MiniBatchGD_withMomentum(X_training, ...
134             Y_training,y_training,GDparams,W,b,lambda,rho);
135         fprintf("Epoch: %d\n",i);
136         epochs_list(i) = i;
137         cost_training_list(i) = ComputeCost(X_training, Y_training, W, b, lambda);
138         accuracy_training_list(i) = ComputeAccuracy(X_training,y_training,W,b);
139         accuracy_validation_list(i) = ComputeAccuracy(X_validation,y_validation,W,b);
140         cost_validation_list(i) = ComputeCost(X_validation, Y_validation, W, b,lambda);
141         GDparams.eta = GDparams.eta * eta_decay_rate;
142     end
143     %Plot cost in training and validation per epoch.
144     title_text = "Cost validation/training with rho = " + num2str(rho);
145     PlotCost(epochs_list, cost_validation_list, cost_training_list,title_text);
146     %Plot accuracy in training and validation per epoch.
147     title_text = "Accuracy validation/training with rho = " + num2str(rho);
148     PlotAccuracy(epochs_list,accuracy_validation_list, accuracy_training_list,title_text);
149     %Final accuracy in validation dataset and training dataset:
150     accuracy_final_training = ComputeAccuracy(X_training,y_training,W,b);
151     fprintf("Final training accuracy: %f %", accuracy_final_training);
152     accuracy_final_validation = ComputeAccuracy(X_validation,y_validation,W,b);
153     fprintf("Final validation accuracy: %f %",accuracy_final_validation);
154 end
155 %% Exercise 4: Training your network
156 % All the experiments will be using all the examples.
157
158 %% 4.1 Find reasonable values for the learning rate:
159
160
161 %Regularization term to small value:
162 lambda = 0.000001;
163 rho = 0.9;
164 eta_values = [0.1,0.01,0.001,0.0001,0.00001,0.000001,0.0000001];
165 GDparams.n_batch = 100;
166 GDparams.n_epochs = 10;
167
168 final_cost_training_list = {};
169 for eta = eta_values
170     GDparams.eta = eta;
171     cost_training_list = zeros(1, GDparams.n_epochs);
172     [W,b] = Initialize_Network(m,d,K);
173     fprintf("Eta value: %f\n",eta);
174     for i = 1:GDparams.n_epochs
175         fprintf("Epoch: %d\n",i);
176         [W,b] = MiniBatchGD_withMomentum(X_training, ...
177             Y_training,y_training,GDparams,W,b,lambda,rho);
178         cost_training_list(i) = ComputeCost(X_training, Y_training,W,b,lambda);

```

```

178     end
179     final_cost_training_list{end+1} = cost_training_list;
180 end
181 epochs_list = zeros(1,GDparams.n_epochs);
182 for i = 1: GDparams.n_epochs
183     epochs_list(i) = i;
184 end
185 %Plot the cost per each case of eta.
186 PlotCost_Eta_values(epochs_list, eta_values, final_cost_training_list);
187
188 %% 4.2 Coarse-search random.
189
190 GDparams.n_batch = 100;
191 GDparams.n_epochs = 10;
192 decay_rate= 0.95;
193 rho=0.9;
194 n_pairs= 70;
195 %Learning rate range:
196 eta_max = 0.1;
197 eta_min = 0.001;
198 %Lambda range:
199 lambda_max = 0.1;
200 lambda_min = 1e-7;
201
202 validation_accuracy_list = zeros(1,n_pairs);
203 eta_values_list = zeros(1,n_pairs);
204 lambda_values_list = zeros(1,n_pairs);
205
206 for j = 1:n_pairs
207     fprintf("\n Pair number: %d \n ", j);
208     eta_exp = log10(eta_min) + (log10(eta_max) - log10(eta_min))*rand(1, 1);
209     eta = 10^eta_exp;
210     lambda_exp = log10(lambda_min) + (log10(lambda_max) - log10(lambda_min))*rand(1, 1);
211     lambda = 10^lambda_exp;
212
213     GDparams.eta = eta;
214     [W,b] = Initialize_Network(m,d,K);
215
216     for i=1: GDparams.n_epochs
217         fprintf("Epoch: %d\n", i);
218         [W, b] = MiniBatchGD_withMomentum(X_training, Y_training,y_training, GDparams, ...
219             W, b, lambda, rho);
219         GDparams.eta = GDparams.eta*decay_rate;
220     end
221     validation_accuracy_list(j) = ComputeAccuracy(X_validation,y_validation,W,b);
222     eta_values_list(j) = GDparams.eta;
223     lambda_values_list(j) = lambda;
224 end
225
226 %Sort the array:
227 [validation_accuracy_list,validation_indexes] = sort(validation_accuracy_list,'descend');
228 eta_values_list = eta_values_list(validation_indexes);
229 lambda_values_list = lambda_values_list(validation_indexes);
230
231 %Write file:
232 filename = 'exercise_4.2_coarse_search.txt';
233 fid = fopen(filename,'wt');
234 fprintf(fid,"\nAccuracy values:\n");
235 fprintf(fid,'%f\t',validation_accuracy_list);
236 fprintf(fid,"\nEta values:\n");
237 fprintf(fid,'%f\t',eta_values_list);
238 fprintf(fid,"\nLambda values:\n");
239 fprintf(fid,'%f\t',lambda_values_list);
240 fprintf(fid,"\n");
241 fclose(fid);

```

```

242
243 %% 4.2 Coarse-fine search random.
244 %Better values for lambda search.
245
246 GDparams.n_batch = 100;
247 GDparams.n_epochs = 10;
248 decay_rate= 0.95;
249 rho=0.9;
250 n_pairs= 70;
251 %Learning rate range:
252 eta_max = 0.07;
253 eta_min = 0.01;
254 %Lambda range:
255 lambda_max = 0.005;
256 lambda_min = 1e-6;
257
258 validation_accuracy_list = zeros(1,n_pairs);
259 eta_values_list = zeros(1,n_pairs);
260 lambda_values_list = zeros(1,n_pairs);
261
262 for j = 1:n_pairs
263     fprintf("\n Pair number: %d \n ", j);
264     eta_exp = log10(eta_min) + (log10(eta_max) - log10(eta_min))*rand(1, 1);
265     eta = 10^eta_exp;
266     lambda_exp = log10(lambda_min) + (log10(lambda_max) - log10(lambda_min))*rand(1, 1);
267     lambda = 10^lambda_exp;
268
269     GDparams.eta = eta;
270     [W,b] = Initialize_Network(m,d,K);
271
272     for i=1: GDparams.n_epochs
273         fprintf("Epoch: %d\n", i);
274         [W, b] = MiniBatchGD_withMomentum(X_training, Y_training,y_training, GDparams, ...
275             W, b, lambda, rho);
276         GDparams.eta = GDparams.eta*decay_rate;
277     end
278     validation_accuracy_list(j) = ComputeAccuracy(X_validation,y_validation,W,b);
279     eta_values_list(j) = GDparams.eta;
280     lambda_values_list(j) = lambda;
281 end
282
283 %Sort the array:
284 [validation_accuracy_list,validation_indexes] = sort(validation_accuracy_list,'descend');
285 eta_values_list = eta_values_list(validation_indexes);
286 lambda_values_list = lambda_values_list(validation_indexes);
287
288 %Write file:
289 filename = 'exercise_4.2_coarse_fine_search.txt';
290 fid = fopen(filename,'wt');
291 fprintf(fid,"\nAccuracy values:\n");
292 fprintf(fid,'%f\t',validation_accuracy_list);
293 fprintf(fid,"\nEta values:\n");
294 fprintf(fid,'%f\t',eta_values_list);
295 fprintf(fid,"\nLambda values:\n");
296 fprintf(fid,'%f\t',lambda_values_list);
297 fprintf(fid,"\n");
298 fclose(fid);
299
300 %% 4.3 Best hyper-parameters for training the network.
301 %Best combination
302
303 m= 50;
304
305 %Train with all the training data:
306 [X_training1, Y_training1, y_training1] = LoadBatch('./data_batch_1.mat');

```



```

306 [X_training2, Y_training2, y_training2] = LoadBatch('./data_batch_2.mat');
307 [X_training3, Y_training3, y_training3] = LoadBatch('./data_batch_3.mat');
308 [X_training4, Y_training4, y_training4] = LoadBatch('./data_batch_4.mat');
309 [X_training5, Y_training5, y_training5] = LoadBatch('./data_batch_5.mat');
310
311 %Test dataset
312 [X_test, Y_test, y_test] = LoadBatch('./test_batch.mat');
313
314 %except 1000 samples for validation: using the validation dataset
315 %(data_batch_2.mat)
316 X_validation = X_training2(:, 1:1000);
317 Y_validation = Y_training2(:, 1:1000);
318 y_validation = y_training2(:, 1:1000);
319 X_training2 = X_training2(:, 1001:10000);
320 Y_training2 = Y_training2(:, 1001:10000);
321 y_training2 = y_training2(:, 1001:10000);
322
323 %All training data together:
324 X_training = [X_training1, X_training2, X_training3];
325 Y_training = [Y_training1, Y_training2, Y_training3];
326 y_training = [y_training1, y_training2, y_training3];
327
328 %Pre-processing:
329 mean_X = mean(X_training, 2);
330 X_training = X_training - repmat(mean_X, [1, size(X_training, 2)]);
331 X_validation = X_validation - repmat(mean_X, [1, size(X_validation, 2)]);
332 X_test = X_test - repmat(mean_X, [1, size(X_test, 2)]);
333
334 %Variables
335 d = size(X_training,1);
336 N = size(X_training,2);
337 K = size(Y_training,1);
338
339 GDparams.n_batch = 100;
340 GDparams.n_epochs = 30;
341 decay_rate= 0.95;
342 rho= 0.9;
343 %Best combination:
344 GDparams.eta = 0.014001;
345 lambda = 0.003347;
346
347 %Plot the training and validation cost after each epoch of training
348 cost_training_list = zeros(1, GDparams.n_epochs);
349 cost_validation_list = zeros(1, GDparams.n_epochs);
350 [W,b] = Initialize_Network(m,d,K);
351 %Original training cost
352 original_training_cost = ComputeCost(X_training, Y_training,W,b,lambda);
353 for i=1:GDparams.n_epochs
354     [W, b] = MiniBatchGD_withMomentum(X_training, Y_training,y_training, GDparams, W, ...
355         b, lambda, rho);
356     cost_training_list(i) = ComputeCost(X_training, Y_training,W,b,lambda);
357     cost_validation_list(i) = ComputeCost(X_validation, Y_validation,W,b,lambda);
358     GDparams.eta = GDparams.eta*decay_rate;
359     if cost_training_list(i) > 3*original_training_cost
360         fprintf("Cost_training(i) > 3*original_training_cost");
361         i = GDparams.n_epochs;
362     end
363 end
364 epochs_list = zeros(1, GDparams.n_epochs);
365 for i=1:GDparams.n_epochs
366     epochs_list(i) = i;
367 end
368
369 %Plot the loss per epochs in validation and in training:

```

```

370 title_text = "Validation/Training cost per epoch";
371 PlotCost(epochs_list, cost_validation_list, cost_training_list,title_text)
372
373 %Final accuracy in testing and training:
374 final_accuracy_test = ComputeAccuracy(X_test,y_test,W,b);
375 fprintf("Final accuracy in test: %f % \n", final_accuracy_test);
376 final_accuracy_training = ComputeAccuracy(X_training,y_training,W,b);
377 fprintf("Final accuracy in training: %f % \n", final_accuracy_training);
378
379 %% Functions implementation
380
381
382 function PlotCost_Eta_values(epochs_list, eta_values, final_cost_training_list)
383     figure;
384     for i = 1:length(eta_values)
385         text_label = "eta = " + num2str(eta_values(i));
386         plot(epochs_list,final_cost_training_list{i},'DisplayName',text_label);
387         hold on;
388     end
389     hold off;
390     xlabel('Epochs');
391     ylabel('Cost');
392     legend;
393 end
394
395 %Plot accuracy for validation/test - training dataset per epoch.
396 function PlotAccuracy(epochs, accuracy_validation, accuracy_training,title_text)
397     figure;
398     plot(epochs,accuracy_validation,epochs,accuracy_training);
399     title(title_text);
400     xlabel('Epochs');
401     ylabel('Accuracy %');
402     legend('Accuracy in validation', 'Accuracy in training','Location','southwest');
403 end
404
405 %Plot the loss/cost for validation/test -training dataset per epoch.
406 function PlotCost(epochs, cost_validation, cost_training,title_text)
407     figure;
408     plot(epochs,cost_validation,epochs,cost_training);
409     title(title_text);
410     xlabel('Epochs');
411     ylabel('Cost');
412     legend('Cost in validation', 'Cost in training','Location','southwest');
413 end
414
415 %Function to initialize the values of the Weight matrices.
416 function [W,b] = Initialize_Network(m,d,K)
417     W1 = randn(m,d)*0.001;
418     W2 = randn(K,m)*0.001;
419     b1 = zeros(m, 1);
420     b2 = zeros(K, 1);
421     W = {W1,W2};
422     b = {b1, b2};
423 end
424
425 function [X,Y,y] = LoadBatch(filename)
426     A = load(filename);
427     X = double(A.data');
428     X= X/255;
429     y = double(A.labels') + 1;
430     vec = ind2vec(y);
431     Y = full(vec); %One-hot representation
432 end
433
434 %Function to compute the cost of the 2-layer network.

```

```

435 function J = ComputeCost(X,Y,W,b,lambda)
436     [P,h] = EvaluateClassifier(X,W,b);
437     crossentropy_term = sum(diag(-log(double(Y)'*P)));
438     reg_term = sum(W{1}(:).^2)+sum(W{2}(:).^2);
439     J = (1/(size(X,2))*crossentropy_term)+(lambda*reg_term);
440 end
441
442 function [P,h] = EvaluateClassifier(X,W,b)
443     s1 = W{1}*X + b{1};
444     h = max(0, s1);
445     s = W{2} * h + b{2};
446     P = exp(s)./sum(exp(s));
447 end
448
449 %Function to compute the accuracy in a 2-layer neural network:
450 function acc = ComputeAccuracy(X,y,W,b)
451     [P,h] = EvaluateClassifier(X, W, b);
452     [~,pred_idx] = max(P);
453     acc = (sum(pred_idx==y)/size(X,2)*100);
454 end
455
456 function [grad_W, grad_b] = ComputeGradients(X, Y, P, h, W,b,lambda)
457     k_x = size(X,2);
458     W1 = cell2mat(W(1));
459     b1 = cell2mat(b(1));
460     W2 = cell2mat(W(2));
461     b2 = cell2mat(b(2));
462     grad_W1 = zeros(size(W1));
463     grad_W2 = zeros(size(W2));
464     grad_b1 = zeros(size(b1));
465     grad_b2 = zeros(size(b2));
466
467     for i= 1:k_x
468         P_i = P(:, i);
469         h_i = h(:, i);
470         Y_i = Y(:, i);
471         X_i = X(:, i);
472         g = -(Y_i-P_i)';
473         grad_b2 = grad_b2 + g';
474         grad_W2 = grad_W2 + g'*h_i';
475
476         h_i(find(h_i > 0)) = 1;
477         g = g*W2*diag(h_i);
478
479         grad_b1 = grad_b1 + g';
480         grad_W1 = grad_W1 + g'*X_i';
481     end
482     grad_b1 = grad_b1/k_x;
483     grad_W1 = grad_W1/k_x + 2*lambda*W1;
484     grad_b2 = grad_b2/k_x;
485     grad_W2 = grad_W2/k_x + 2*lambda*W2;
486
487     grad_b = {grad_b1, grad_b2};
488     grad_W = {grad_W1, grad_W2};
489 end
490
491 function [grad_W, grad_b] = ComputeGradsNumSlow(X, Y, W, b, lambda, h)
492     grad_W = cell(numel(W), 1);
493     grad_b = cell(numel(b), 1);
494     for j=1:length(b)
495         grad_b{j} = zeros(size(b{j}));
496         for i=1:length(b{j})
497             b_try = b;
498             b_try{j}(i) = b_try{j}(i) - h;
499             c1 = ComputeCost(X, Y, W, b_try, lambda);

```

```

500         b_try = b;
501         b_try{j}(i) = b_try{j}(i) + h;
502         c2 = ComputeCost(X, Y, W, b_try, lambda);
503         grad_b{j}(i) = (c2-c1) / (2*h);
504     end
505 end
506 for j=1:length(W)
507     grad_W{j} = zeros(size(W{j}));
508     for i=1:numel(W{j})
509         W_try = W;
510         W_try{j}(i) = W_try{j}(i) - h;
511         c1 = ComputeCost(X, Y, W_try, b, lambda);
512         W_try = W;
513         W_try{j}(i) = W_try{j}(i) + h;
514         c2 = ComputeCost(X, Y, W_try, b, lambda);
515         grad_W{j}(i) = (c2-c1) / (2*h);
516     end
517 end
518 end
519
520
521 %Mini batch function implementation without momentum and learning rate
522 %decay.
523 function [Wstar, bstar] = MiniBatchGD(X,Y,y,GDparams,W,b,lambda)
524     N = size(X,2);
525     eta = GDparams.eta;
526     n_batch = GDparams.n_batch;
527
528     for j=1:N/n_batch
529         j_start = (j-1)*n_batch + 1;
530         j_end = j*n_batch;
531         indx = j_start:j_end;
532         Xbatch = X(:,indx);
533         Ybatch = Y(:,indx);
534         [P,h] = EvaluateClassifier(Xbatch,W,b);
535         [grad_W, grad_b] = ComputeGradients(Xbatch,Ybatch,P,h, W,b,lambda);
536         W{1} = W{1} - eta*grad_W{1};
537         W{2} = W{2} - eta*grad_W{2};
538         b{1} = b{1} - eta*grad_b{1};
539         b{2} = b{2} - eta*grad_b{2};
540     end
541     bstar = b;
542     Wstar = W;
543 end
544
545 function [v_W, v_b] = InitializeMomentum(W,b)
546     v_b = {zeros(size(b{1})), zeros(size(b{2}))};
547     v_W = {zeros(size(W{1})), zeros(size(W{2}))};
548 end
549
550 function [Wstar, bstar] = MiniBatchGD_withMomentum(X, Y,y,GDparams,W,b,lambda,rho)
551     N = size(X,2);
552     eta = GDparams.eta;
553     n_batch = GDparams.n_batch;
554     [v_W,v_b] = InitializeMomentum(W,b);
555
556     for j=1:N/n_batch
557         j_start = (j-1)*n_batch + 1;
558         j_end = j*n_batch;
559         indx = j_start:j_end;
560         Xbatch = X(:,indx);
561         Ybatch = Y(:,indx);
562         [P,h] = EvaluateClassifier(Xbatch,W,b);
563         [grad_W, grad_b] = ComputeGradients(Xbatch,Ybatch,P,h, W,b,lambda);
564         %Update with momentum:

```

```

565         v_W{1} = rho*v_W{1} + eta*grad_W{1};
566         v_b{1} = rho*v_b{1} + eta*grad_b{1};
567         v_W{2} = rho*v_W{2} + eta*grad_W{2};
568         v_b{2} = rho*v_b{2} + eta*grad_b{2};
569         W{1} = W{1} - v_W{1};
570         b{1} = b{1} - v_b{1};
571         W{2} = W{2} - v_W{2};
572         b{2} = b{2} - v_b{2};
573     end
574     bstar = b;
575     Wstar = W;
576 end

```